



PSI

Center for Scientific Computing,
Theory and Data

How to write your workflow

Lessons I learned writing the koopmans package

Edward Linscott

LMS Seminar, 12 March 2025

A brief history of the koopmans code



A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE
- the first workflow was Koopmans Δ SCF for molecules; others quickly followed

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE
- the first workflow was Koopmans Δ SCF for molecules; others quickly followed
- reusability of workflows as subworkflows became important (e.g. Wannierisation)

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE
- the first workflow was Koopmans Δ SCF for molecules; others quickly followed
- reusability of workflows as subworkflows became important (e.g. Wannierisation)
- start to add tests, documentation, etc, ...

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE
- the first workflow was Koopmans Δ SCF for molecules; others quickly followed
- reusability of workflows as subworkflows became important (e.g. Wannierisation)
- start to add tests, documentation, etc, ...

koopmans

A brief history of the koopmans code

- I arrived in THEOS in 2019, started learning how to run Koopmans functionals – and it was painful!
- I was familiar with ASE so wrote some simple scripts
- decided this was something we wanted for everyone
- ... and having been advised not to use AiiDA we stuck with the ASE approach
- implemented kcp.x support in ASE
- the first workflow was Koopmans Δ SCF for molecules; others quickly followed
- reusability of workflows as subworkflows became important (e.g. Wannierisation)
- start to add tests, documentation, etc, ...



- (more recently) substantial changes to integrate with AiiDA (I'll discuss this later)

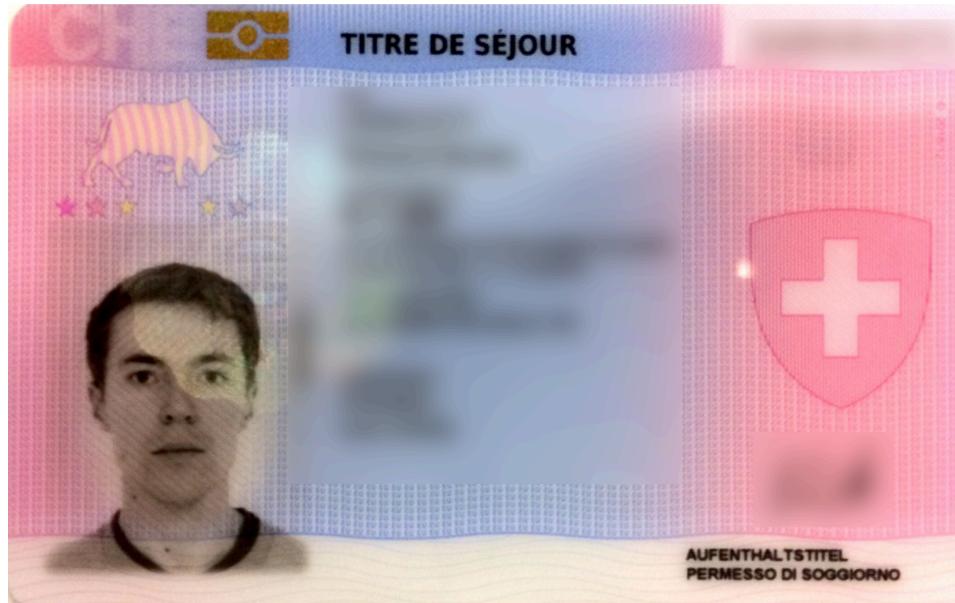
**.... but what have I learned
throughout this process?**



A scene from The Simpsons. Homer Simpson, wearing a tan shirt and grey pants, is standing on the left, gesturing with his hands while speaking. Bart Simpson, in his orange shirt, is looking up at him. In the foreground, the backs of two other characters are visible: a black-haired person in a blue shirt and a blonde person in a white shirt. The background shows a green hillside with a blue sky and clouds.

... not quite!

2019



2024



What I learned when interfacing koopmans with AiiDA

koopmans

Simple by design



Simple by design

- local execution only



Simple by design

- local execution only
- serial step execution (even when steps are independent!)



Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files



Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

- remote execution

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

- remote execution
- parallel step execution

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

- remote execution
- parallel step execution
- outputs stored in a database

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

- remote execution
- parallel step execution
- outputs stored in a database
- installation more involved

Simple by design

- local execution only
- serial step execution (even when steps are independent!)
- direct access to input/output files
- simple installation

Powerful by design

- remote execution
- parallel step execution
- outputs stored in a database
- installation more involved

We could really benefit from a lot of these features

Goals

- be able to execute in parallel and remotely with AiiDA

Goals

- be able to execute in parallel and remotely with AiiDA
- UI should be as similar as possible

Goals

- be able to execute in parallel and remotely with AiiDA
- UI should be as similar as possible
- old mode of running koopmans should still work

Goals

- be able to execute in parallel and remotely with AiiDA
- UI should be as similar as possible
- old mode of running koopmans should still work
- minimal/no duplication of logic

Refactoring

A lot of “bad” design in koopmans needed refactoring

Refactoring

A lot of “bad” design in koopmans needed refactoring

- abstraction of many operations e.g.

```
with open(filename, 'r') as f:  
    data = f.read()
```

Refactoring

A lot of “bad” design in koopmans needed refactoring

- abstraction of many operations e.g.

```
with open(filename, 'r') as f:  
    data = f.read()
```

becomes

```
self.engine.read(filename)
```

Refactoring

A lot of “bad” design in koopmans needed refactoring

- abstraction of many operations e.g.

```
with open(filename, 'r') as f:  
    data = f.read()
```

becomes

```
self.engine.read(filename)
```

- ... and, more generally, many responsibilities are moved to the engine (reading/writing files, running calculations, checking the status of calculations, loading pseudopotentials, etc.)

Refactoring

- making everything more “pure”

Refactoring

- making everything more “pure” e.g. removing all reliance on shared directories

Refactoring

- making everything more “pure” e.g. removing all reliance on shared directories

```
calc_nscf.parameters.outdir = calc_scf.parameters.outdir
```

Refactoring

- making everything more “pure” e.g. removing all reliance on shared directories

```
calc_nscf.parameters.outdir = calc_scf.parameters.outdir
```

becomes

```
calc_nscf.link(calc_scf.parameters.outdir, 'tmp')
```

where

```
def link(self, src, dst):  
    self.engine.link(src=src, dst=self/dst)
```

Note: none of this refactoring
is specific to AiiDA!

At the AiiDA end



At the AiiDA end

- conversion of calculators from ASE to AiiDA and back

At the AiiDA end

- conversion of calculators from ASE to AiiDA and back
- verdi presto for simplified AiiDA setup

At the AiiDA end

- conversion of calculators from ASE to AiiDA and back
- verdi presto for simplified AiiDA setup
- verdi dump for dumping AiiDA database to a local file structure

**Writing workflows well is
hard...**

**Writing workflows well is
hard... how can it best be
done?**

Common Workflow Language

Basic Concepts of CWL



COMMON
WORKFLOW
LANGUAGE

Basic Concepts of CWL



- “an open standard for describing how to run command line tools and connect them to create workflows”

Basic Concepts of CWL



- “an open standard for describing how to run command line tools and connect them to create workflows”
- separate runners are required to execute the workflows; a CWL workflow only contains the logic of the workflows

Basic Concepts of CWL



- “an open standard for describing how to run command line tools and connect them to create workflows”
- separate runners are required to execute the workflows; a CWL workflow only contains the logic of the workflows
- introduced in 2014; version 1.2 released in 2020

Basic Concepts of CWL



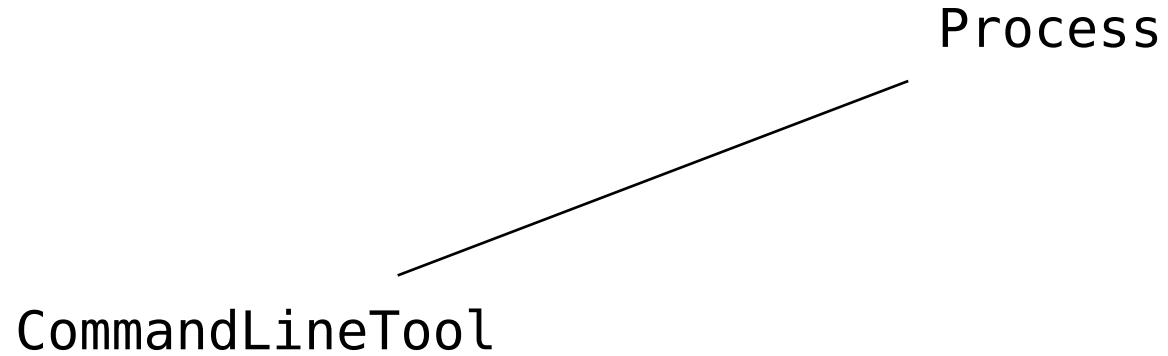
- “an open standard for describing how to run command line tools and connect them to create workflows”
- separate runners are required to execute the workflows; a CWL workflow only contains the logic of the workflows
- introduced in 2014; version 1.2 released in 2020
- mostly used by bioinformatics community

Basic Concepts of CWL

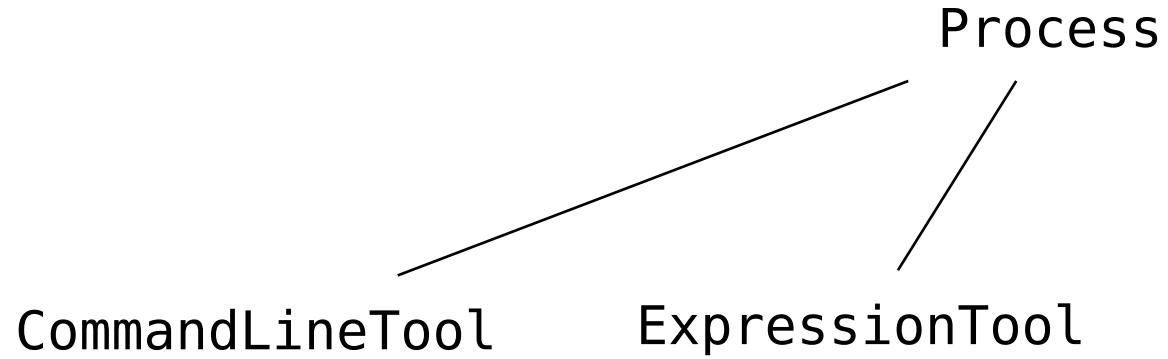


Process

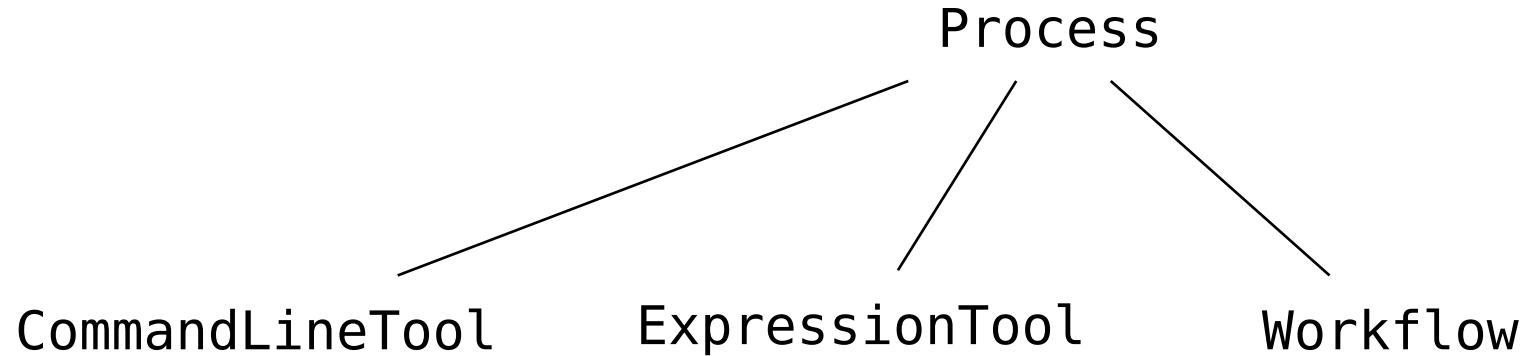
Basic Concepts of CWL



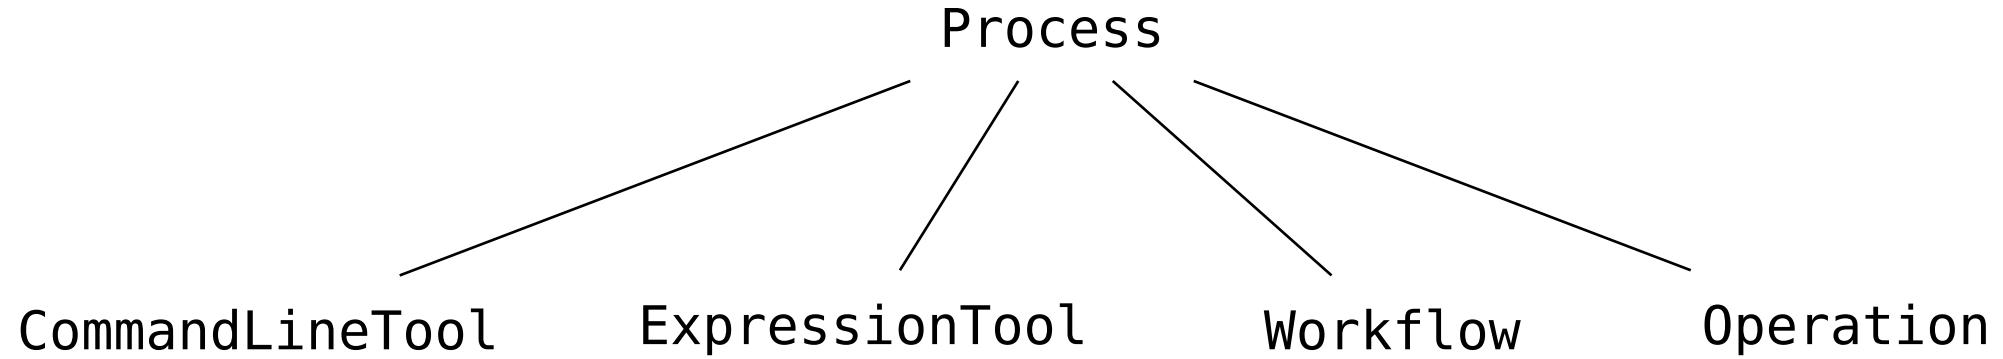
Basic Concepts of CWL



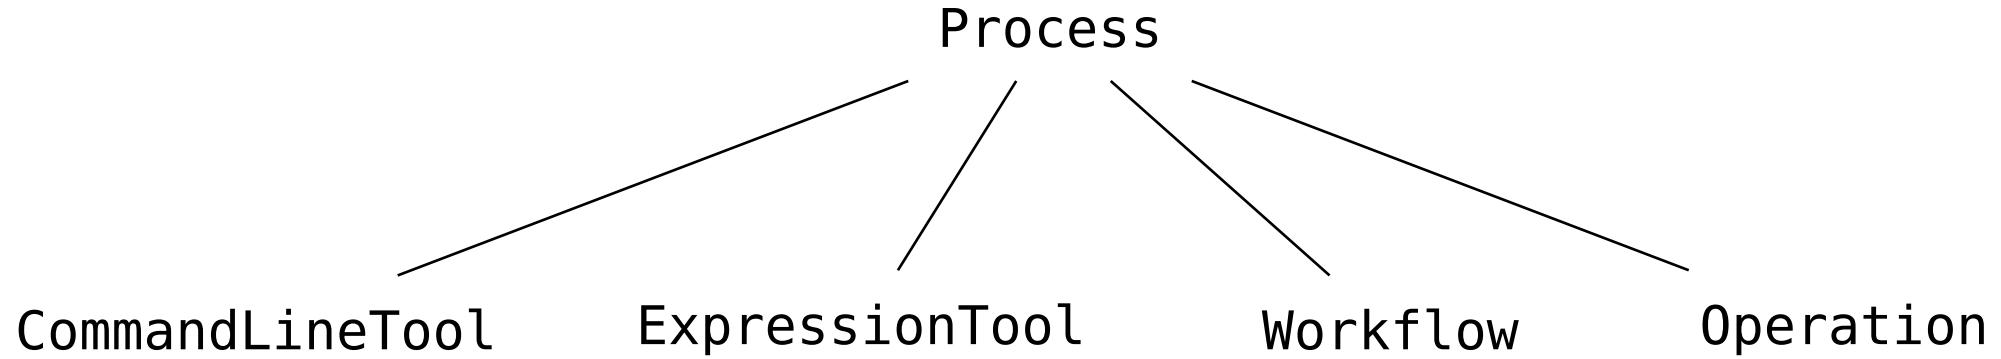
Basic Concepts of CWL



Basic Concepts of CWL



Basic Concepts of CWL



Every Process has inputs and outputs

CommandLineTool

```
echo.cwl

cwlVersion: v1.2
class: CommandLineTool

inputs:
  message:
    type: string
    default: "Hello World"
    inputBinding:
      position: 1
outputs: []

baseCommand: echo
```

CommandLineTool

```
echo.cwl

cwlVersion: v1.2
class: CommandLineTool

inputs:
  message:
    type: string
    default: "Hello World"
    inputBinding:
      position: 1
outputs: []

baseCommand: echo
```

Called via \$ cwltool echo.cwl

ExpressionTool

```
uppercase.cwl
```

```
cwlVersion: v1.2
class: ExpressionTool
requirements:
  InlineJavascriptRequirement: {}

inputs:
  message: string
outputs:
  uppercase_message: string

expression: |
  ${ return {"uppercase_message": inputs.message.toUpperCase()}; }
```

ExpressionTool



uppercase.cwl

```
cwlVersion: v1.2
class: ExpressionTool
requirements:
  InlineJavascriptRequirement: {}

inputs:
  message: string
outputs:
  uppercase_message: string

expression: |
  ${ return {"uppercase_message": inputs.message.toUpperCase()}; }
```

Called via \$ cwltool uppercase.cwl --message='goes to 11'

Workflow

```
echo_uppercase.cwl

cwlVersion: v1.2
class: Workflow

requirements:
  InlineJavascriptRequirement: {}

inputs:
  message: string
outputs:
  out:
    type: string
    outputSource: uppercase/uppercase_message
```

Workflow



```
steps:  
  echo:  
    run: echo.cwl  
    in:  
      message: message  
    out: [out]  
  uppercase:  
    run: uppercase.cwl  
    in:  
      message:  
        source: echo/out  
    out: [uppercase_message]
```

Workflow

```
echo_uppercase.cwl

cwlVersion: v1.2
class: Workflow

requirements:
  InlineJavascriptRequirement: {}

inputs:
  message: string
outputs:
  out:
    type: string
    outputSource: uppercase/uppercase_message
```

Workflow



```
steps:  
  echo:  
    run: echo.cwl  
    in:  
      message: message  
    out: [out]  
  uppercase:  
    run: uppercase.cwl  
    in:  
      message:  
        source: echo/out  
    out: [uppercase_message]
```

Called via \$ cwltool echo_uppercase.cwl input.json

Operation

```
p_vs_np.cwl
```

```
cwlVersion: v1.2
```

```
class: Operation
```

```
inputs: []
```

```
outputs:
```

```
  result: bool
```

A less silly Operation

```
run_pw.cwl
cwlVersion: v1.2
class: Operation

inputs:
    input_file: File
    pseudopotentials:
        type: array
        items:
            type: File
outputs:
    output_file: File
```

Pros and Cons

- pros:

Pros and Cons

- pros:
 - explicit

Pros and Cons

- pros:
 - explicit
 - composable

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate
 - complicated workflows are **very** complicated to write in CWL (e.g. while)

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate
 - complicated workflows are **very** complicated to write in CWL (e.g. while)
 - ExpressionTool restricted to Javascript

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate
 - complicated workflows are **very** complicated to write in CWL (e.g. while)
 - ExpressionTool restricted to Javascript
 - need to define custom types (see e.g. OPTIMADE, PREMISE)

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate
 - complicated workflows are **very** complicated to write in CWL (e.g. while)
 - ExpressionTool restricted to Javascript
 - need to define custom types (see e.g. OPTIMADE, PREMISE)
 - custom types do not permit defaults

Pros and Cons

- pros:
 - explicit
 - composable
 - customisable
- cons:
 - verbose; lots of boilerplate
 - complicated workflows are **very** complicated to write in CWL (e.g. while)
 - ExpressionTool restricted to Javascript
 - need to define custom types (see e.g. OPTIMADE, PREMISE)
 - custom types do not permit defaults
 - rigorous schemas require willingness from the community

... willingness from the community?



The screenshot shows a GitLab issue page for a project named 'q-e'. The issue, titled 'Suggestions to make .def files more informative and consistent', was created by Edward Linscott 9 months ago and is currently open. The description discusses shortcomings in the .def files and suggests solutions. The sidebar on the left shows other sections like 'Issues' (113) and 'Merge requests' (14). The right side of the page contains fields for 'Assignees' (None), 'Epic' (locked, upgrade plan), 'Labels' (None), 'Milestone' (None), and 'Weight' (None).

Suggestions to make .def files more informative and consistent

Open Issue created 9 months ago by Edward Linscott

Currently, the .def files that define the Quantum ESPRESSO input format have a couple of shortcomings. Below, I discuss these shortcomings and suggest some solutions.

Dimensionality and units

The dimensionality and units of each variable is not well-documented. Yes, the header specifies that unless otherwise specified quantities are in Rydberg atomic units, but the way that alternative units are specified within the info tag is ad-hoc e.g.

```
var sci_vb -type REAL {  
    default { 0 }  
    info {
```

... willingness from the community?

The screenshot shows a Jira interface. On the left is a sidebar with navigation links: Merge requests (14), Manage, Plan, Issues (113), Issue boards, Milestones, Wiki, Requirements, Code, Build, Deploy, and Operate. The Issues link is highlighted. The main content area shows a pull request with the title:

```
val INDEX_SECONDS -type REAL 1  
    default { 1.0+7, or 150 days, i.e. no time limit }
```

A comment below the code reads:

I would propose the introduction of either (a) a new field (`default_explanation` or similar) or (b) a sub-field of `default`. This new field would contain any explanation of the default is contained, while `default` would be strictly reserved for the actual numerical value of the default.

Below this is a section titled "Why should I care?". A tooltip from Julian Geiger says:

Julian Geiger reacted with :thumbsup:

Below the tooltip are three buttons: a thumbs up icon with the number 1, a thumbs down icon with the number 0, and a neutral face icon.

The Jira interface also includes sections for Child items (0), a summary note about child items, and various metadata fields on the right side:

- Labels: None
- Milestone: None
- Weight: None
- Due date: None
- Time tracking: No estimate or time spent
- Health status: None

Building ideas from CWL into koopmans

Building ideas from CWL into koopmans

koopmans is slowly being refactored into “CWL-inspired” Python

```
InputModel = TypeVar('InputModel', bound=BaseModel)
OutputModel = TypeVar('OutputModel', bound=BaseModel)

class Process(ABC, Generic[InputModel, OutputModel]):

    input_model: Type[InputModel]
    output_model: Type[OutputModel]

    def __init__(self, name: str | None = None, **kwargs):
        self.inputs: InputModel = self.input_model(**kwargs)
        self.outputs: OutputModel | None = None
        self.directory: Path | None = None

    def run(self):
        assert self.directory is not None, 'Process directory must be set before running'
        with utils.chdir(self.directory):
            self.dump_inputs()
            self._run()
            assert self.outputs is not None, 'Process outputs must be set when running'
            self.dump_outputs()

    @abstractmethod
    def _run(self):
        ...
```

A simple CommandLineTool

```

class Bin2XMLInput(IOModel):
    binary: File
    class Config:
        arbitrary_types_allowed = True

class Bin2XMLOutput(IOModel):
    xml: File
    class Config:
        arbitrary_types_allowed = True

class Bin2XMLPCommandLineTool(CommandLineTool[Bin2XMLInput, Bin2XMLOutput]):

    input_model = Bin2XMLInput
    output_model = Bin2XMLOutput

    def _pre_run(self):
        super()._pre_run()
        if not self.inputs.binary.exists():
            raise FileNotFoundError(f'{self.inputs.binary} does not exist')

        # Link the input binary file to the directory of this process as input.dat
        dst = self / "input.dat"
        dst.symlink_to(self.inputs.binary)

    @property
    def command(self):
        return Command(executable='bin2xml.x', suffix='input.dat output.xml')

    def _set_outputs(self):
        self.outputs = self.output_model(xml=self / "output.xml")

```

An example composite workflow

```
# Read and run the Koopmans workflow
wf = read('si.json')
wf.run_while()

# Merge the separate occ + emp projections into a set of projections that combines occ + empty
combined_proj_list = [p for block in wf.projections for p in block.projections]
combined_proj_obj = ProjectionBlocks.fromlist([combined_proj_list], ['up'], wf.atoms)

# Find the kcw.x output directory that contains the Hamiltonian to Wannierize
[kcw_ham_calc] = [c for c in wf.calculations if isinstance(c, KoopmansHamCalculator)]
kcw_outdir = File(kcw_ham_calc, kcw_ham_calc.parameters.outdir)

# Construct a Wannierize workflow for the joint occ + empty manifold
wann_wf = WannierizeBlockWorkflow.from_other(wf,
    block=combined_proj_obj[0],
    pw_outdir=kcw_outdir,
    write_tb=True,
    calculate_bands=True)
wann_wf.calculator_parameters['pw2wannier'].prefix = 'kc_kcw'
wann_wf.directory /= '02-joint-wannierize'

# Run the Wannierization
wann_wf.run_while()
```

So what?

koopmans + AiIDA

UI practically unchanged:

```
$ koopmans tio2.json
```

koopmans + AiiDA

UI practically unchanged:

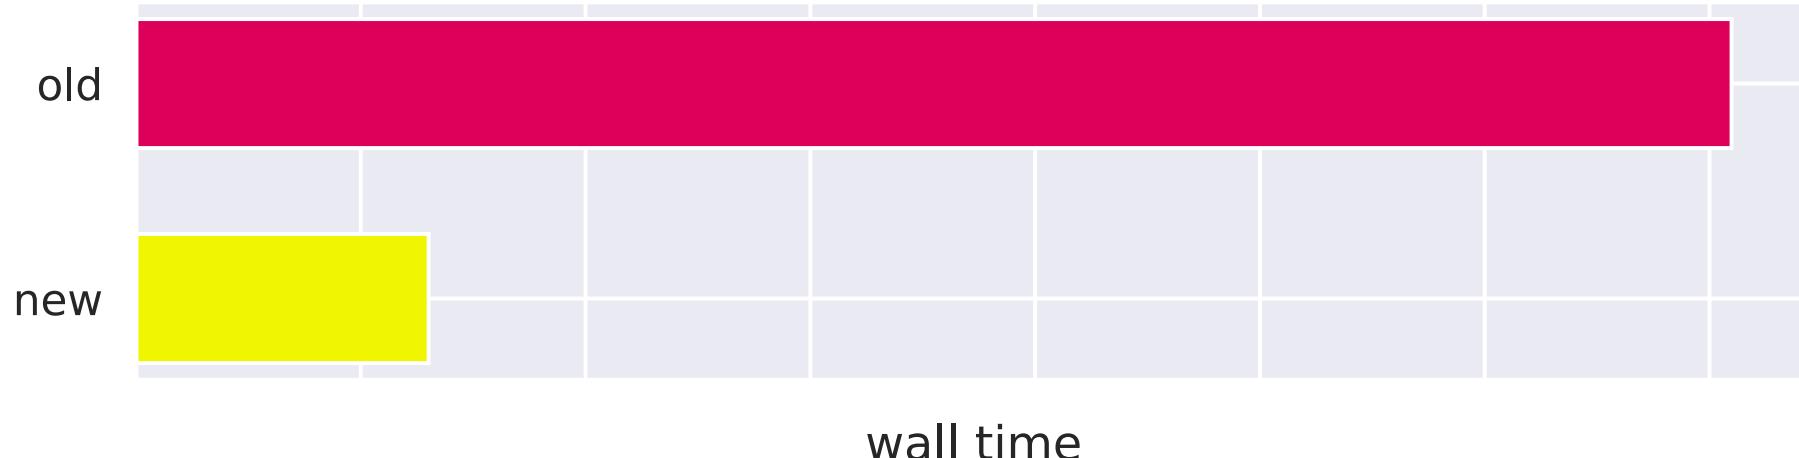
```
$ koopmans tio2.json → $ koopmans run --engine=aiida tio2.json
```

koopmans + AiiDA

UI practically unchanged:

```
$ koopmans tio2.json → $ koopmans run --engine=aiida tio2.json
```

but executed remotely and in parallel:



Outlook

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* **7**, 1–12 (2021)

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

Can we write runner-agnostic workflows?

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* 7, 1–12 (2021)

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

Can we write runner-agnostic workflows?

- yes! CWL the most rigorous way¹, but koopmans (kind of) achieves this

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* 7, 1–12 (2021)

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

Can we write runner-agnostic workflows?

- yes! CWL the most rigorous way¹, but koopmans (kind of) achieves this

Do we want to write runner-agnostic workflows?

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* 7, 1–12 (2021)

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

Can we write runner-agnostic workflows?

- yes! CWL the most rigorous way¹, but koopmans (kind of) achieves this

Do we want to write runner-agnostic workflows?

- probably not...

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* 7, 1–12 (2021)

Workflow-runner-agnostic workflows?

A lot of the pain in koopmans + AiiDA derives from wanting to have two workflow runners:

- localhost (the original koopmans implementation)
- AiiDA

Can we write runner-agnostic workflows?

- yes! CWL the most rigorous way¹, but koopmans (kind of) achieves this

Do we want to write runner-agnostic workflows?

- probably not...

(cf. Do we want to write calculator-agnostic workflows?

- yes! See Common Workflows²)

¹but still requires schemas, so not truly general

²S. P. Huber *et al.* *npj Comput Mater* 7, 1–12 (2021)

If not, what should we do?

What do people want?

If not, what should we do?

What do people want?

- quick to write

If not, what should we do?

What do people want?

- quick to write
- understandable

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!
- be modular wherever possible (e.g. separate responsibilities)

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!
- be modular wherever possible (e.g. separate responsibilities)
- be abstract wherever possible

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!
- be modular wherever possible (e.g. separate responsibilities)
- be abstract wherever possible
- make it easy for people to tweak and combine workflows

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!
- be modular wherever possible (e.g. separate responsibilities)
- be abstract wherever possible
- make it easy for people to tweak and combine workflows
 - (should AiiDA be able to read and dump .cwl files?)

If not, what should we do?

What do people want?

- quick to write
- understandable
- customisable

What should we do when writing code?

- always have inputs and outputs rigorously defined – think functionally!
- be modular wherever possible (e.g. separate responsibilities)
- be abstract wherever possible
- make it easy for people to tweak and combine workflows
 - (should AiiDA be able to read and dump .cwl files?)
- for AiiDA , continue efforts to simplify

Fin

References



S. P. Huber *et al.* Common workflows for computing material properties using different quantum engines. *npj Comput Mater* **7**, 1–12 (2021).